

Lec 22. What next?

Eunjung Kim

UNLIKELY TO BE P-TIME, WHAT NEXT?

POPULAR ALGORITHMIC APPROACHES

- Exponential running time, but as fast as possible.
~> **Exponential-time algorithm**.
- Not exact/optimal solution, but as close to optimal as possible.
~> **Approximation algorithm**. runs in P-time.
- Exponential running time, but exponential w.r.t a small parameter k .
~> **Fixed-parameter algorithm**. runs in time $f(k) \cdot n^{O(1)}$.
- Produced intended solution with high probability.
~> **Randomized algorithm**. can be optimal or approximate.

ALGORITHMS BEYOND THE CLASSIC SET-UP

POPULAR ALGORITHMIC APPROACHES

- **online algorithm**: competitive ratio, memory, update time
- **distributed algorithm**: communication, data size on each processor

DIRECTED HAMILTONIAN PATH

A Hamiltonian path of a directed graph $G = (V, E)$ is a directed path which visits every vertex of G precisely once.

PROBLEM DIRECTED HAMILTONIAN PATH

INPUT a graph G .

QUESTION does G have a Hamiltonian path?

DYNAMIC PROGRAMMING FOR HAM

Naive algorithm: guess all possible tours, $n!$ many of them.

Dynamic programming algorithm in $2^n \cdot n^{O(1)}$ time.

For every vertex subset X and $v \in X$,
 $P[X, v] = 1$ if and only if there exists a Hamiltonian path of $G[X]$ ending in v .

- 1 Initialize: $P[\{w\}, w] = 1$ for every $w \in V$.
- 2 Inductive step: assume $P[X, v]$ is known for all $|X|$ of size i and $v \in X$.
Let's compute $P[X, v]$ for each $|X|$ of size $i + 1$ and $v \in X$.

$$P[X, v] = \bigvee_{w \in X \setminus v} P[X \setminus v, w] \cdot [(w, v) \in E(G)].$$

- 3 There exists a Hamiltonian path of G if and only if there is a vertex v s.t.
 $P[V, v] = 1$.

VERTEX COVER

A vertex cover of a graph is a vertex subset X of G such that X takes at least one endpoint of every edge in G .

VERTEX COVER

INPUT a graph G and an integer k .

QUESTION does G have a vertex cover of size at most k ?

VERTEX COVER is NP-complete; reduction from INDEPENDENT SET.

APPROXIMATION FOR VERTEX COVER

GREEDY ALGORITHM FOR VERTEX COVER

- 1 Greedily find a maximal matching M .
- 2 Output $V(M)$.

- The output of the algorithm is indeed a vertex cover. (Why?)
- Analysis:

$$|M| \leq \text{opt vc} \leq \text{output vc} = 2 \cdot |M|.$$

Therefore, $\text{output vc} \leq 2 \cdot \text{opt vc}$.

- 2-approximation: algorithm outputs a solution no larger than twice the optimal.

PARAMETERIZED VERTEX COVER

PARAMETERIZED VERTEX COVER

INPUT a graph G and an integer k .

PARAMETER k .

QUESTION does G have a vertex cover of size at most k ?

FPT-ALGORITHM FOR VERTEX COVER

Algorithm **VC**(G, k)

- 1 If G has no edge **return** YES.
- 2 Else if $k = 0$ **return** NO.
- 3 Else (comment: G has an edge and $k > 0$)
 - 1 Pick an edge uv .
 - 2 Return **VC**($G - u, k - 1$) or **VC**($G - v, k - 1$).

FPT-ALGORITHM FOR VERTEX COVER

Runtime analysis:

- Recursive calls of $\mathbf{VC}(G, k)$ can be expressed as a branching tree \mathcal{T} .
- Parameter k strictly decreases by depth \rightsquigarrow depth at most k .
- \mathcal{T} has at most 2^k leaves \rightsquigarrow runs in time $2^k \cdot \text{poly}(n)$.

RANDOMIZED APPROXIMATION FOR MAX CUT

MAX CUT

INPUT a graph $G = (V, E)$.

QUESTION find a bipartition of V into V_1, V_2 maximizing the crossing edges between V_1 and V_2 .

MAX CUT is NP-complete; reduction from 3SAT via a series of reductions.

RANDOMIZED APPROXIMATION FOR MAX CUT

RANDOM GREEDY ALGORITHM FOR MAX CUT

- 1 For each vertex v , decide where to put v with prob 0.5, independently at random.
 - 2 Count the crossing edges between V_1 and V_2 .
- For arbitrary edge e , probability that e counts as 'crossing' is $1/2$.
 - $E(\# \text{crossing edges}) = \sum_{e \in E} P(e \text{ is crossing}) = |E|/2$.
 - In **expectation**, the output is at least half the optimal solution.
 - Repeat the greedy algorithm several times \rightsquigarrow turn the 'expectation' into probability guarantee.
 - Derandomization techniques.

ONLINE ALGORITHM FOR SECRETARY PROBLEM

SECRETARY PROBLEM

- 1 Input: n candidates for a secretary job comes for an interview, in a random order.
- 2 Goal: find the best candidate
- 3 Set-up: each secretary i has competence $w(i)$. Once the interview is done, you have to decide **immediately** to hire the candidate or not. A candidate you once rejected can never be recalled. You know the number of candidates.

Competitive ratio:

$w(\text{the chosen candidate}) / w(\text{the best one among } n \text{ candidates})$

ONLINE ALGORITHM FOR SECRETARY PROBLEM

Algorithm

- 1 Reject the first n/e candidates.
- 2 Afterwards, hire the first candidate who is as good as the best one among the first n/e .

TWO-PARTY COMMUNICATION: EQUALITY

EQUALITY

- Alice has $x \in \{0, 1\}^n$, Bob has $y \in \{0, 1\}^n$.
- They want to decide if $x = y$ with minimum number of bit exchanges.

TWO-PARTY COMMUNICATION: EQUALITY

Deterministic communication

- 1 Alice sends to Bob x ; n -bits.
- 2 Bob compares x and y .

Randomized communication: public randomness

- 1 Uniform random $z \in \{0, 1\}^n$; both Alice and Bob can access it.
- 2 Alice computes $x \cdot z$ and sends the outcome to Bob; 1-bits.
- 3 Bob computes $y \cdot z$ and reject if $y \cdot z \neq x \cdot z$.

TWO-PARTY COMMUNICATION: EQUALITY

Deterministic communication

- 1 Alice sends to Bob x ; n -bits.
- 2 Bob compares x and y .

Randomized communication: public randomness

- 1 Uniform random $z \in \{0, 1\}^n$; both Alice and Bob can access it.
- 2 Alice computes $x \cdot z$ and sends the outcome to Bob; 1-bits.
- 3 Bob computes $y \cdot z$ and reject if $y \cdot z \neq x \cdot z$.

DISTRIBUTED COMPUTATION: LEADER ELECTION

LEADER ELECTION IN DISTRIBUTED COMPUTING

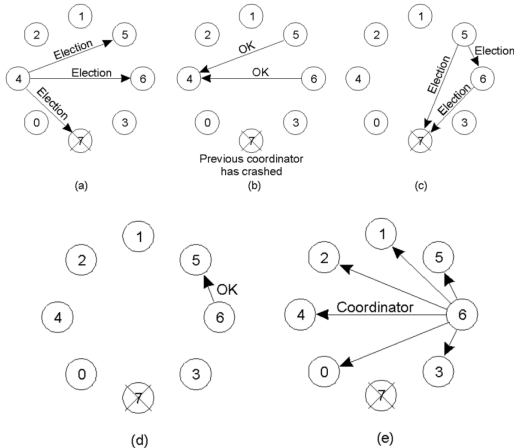
- n nodes (processors) with distinct ID, every processor knows the IP address of every other processor and their IDs.
- Some processors may crash at times.
- Want to make sure that the highest ID node currently alive has the token.
- Protocol so that all processors know the leader (with little communication overhead).

DISTRIBUTED COMPUTATION: LEADER ELECTION

Bully Algorithm

- 1 Three types of messages: (i) Election started, (ii) I'm alive, (iii) I'm the leader.
- 2 If a processor with the highest ID wakes up (after crash), it sends "I'm the leader" message.
- 3 If a node detects a failure of the leader, it sends "Election" to all higher ID nodes.
- 4 If a node receives "Election" message, it sends "I'm alive" to the sender and sends "Election" message to all higher ID nodes.
- 5 If a node does not receive "I'm alive" message after sending "Election", it sends "I'm the leader" message to everyone.
- 6 If a node receives "I'm the leader" message, it considers the sender the leader.

DISTRIBUTED COMPUTATION: LEADER ELECTION



DISTRIBUTED COMPUTATION: COMPACT LOCAL CERTIFICATE

CERTIFICATE

- You're a processor (node) in a network (graph) with ID of length $\log n$. You are only aware of your neighbors and nothing beyond it.
- An external coordinator can give each node v some data $\ell(v)$.
- Each node can compute whatever they likes (unlimited computing power).
- One round of communication: each node communicates with its neighbors some message.
- After-communication computation: each node, based on what it heard from its neighbors, makes a decision and say "yes" or "no".

DISTRIBUTED COMPUTATION: COMPACT LOCAL CERTIFICATE

Can the data $\ell(v)$ be designed (as small size as possible) so that

- All nodes say "yes" if property P holds.
- Some node says "no" if property P does not hold.

$O(1)$ -BIT LOCAL CERTIFICATE FOR BIPARTITENESS

Can the nodes collectively decide if the underlying network is bipartite?

$O(\log n)$ -BIT LOCAL CERTIFICATE FOR TREE

Can the nodes collectively decide if the underlying network is a forest?